Learning Instance-Dependent Parameters for Fixed-Depth QAOA via Graph Neural Networks

Molena Huynh
(Dated: April 9, 2026)

Parameter optimization remains a central bottleneck in fixed-depth QAOA because each new graph instance typically requires another classical outer-loop search. We demonstrate that learning QAOA parameters via graph neural networks improves optimization performance and generalization across held-out graph instances in the fixed-depth regime studied here. Our setting is weighted MaxCut on transcriptomic co-expression graphs, where a depth-2 graph neural initializer is compared against random or heuristic initialization, an affine graph-feature regressor, and direct multi-start classical search. On six held-out graphs, the learned initializer attains mean approximation ratio 0.8682 versus 0.8686 for direct search while reducing median online latency from 936.1255 ms to 0.3652 ms. Under local depolarizing noise rates 0.01, 0.02, and 0.05, it retains mean approximation ratios 0.8433, 0.8268, and 0.7850, remaining close to the classical-search angles on the same graphs. Beyond the benchmark, we provide an exact QUBO/Ising encoding for a cohort-level biomedical objective, a second-order regret bound linking angle error to objective loss, an expressivity separation from affine graph-feature regression, a finite-family non-plateau statement for the fixed-depth regime, and explicit resource scaling. The claim is therefore not quantum advantage. It is that graph-conditioned parameter prediction can be analyzed formally and can replace repeated parameter search for a structured weighted graph family while preserving nearly all of the downstream objective value in both ideal and noisy simulation.

I. INTRODUCTION

Variational quantum algorithms are often limited less by ansatz definition than by parameter selection. In QAOA, fixed depth p leaves only a small number of angles to optimize, but those angles must ordinarily be searched separately for every new graph instance [1–9]. If the same structured family of graphs appears repeatedly, this repeated outer-loop optimization becomes a natural target for learned parameter prediction.

This paper studies that setting for weighted MaxCut instances derived from transcriptomic co-expression graphs. A graph neural network predicts the depth-2 QAOA parameters $(\gamma_1, \gamma_2, \beta_1, \beta_2)$ from the input graph, and those predicted angles are evaluated by exact statevector simulation against direct multi-start classical search. The empirical question is narrow: can graph-conditioned prediction replace repeated per-instance parameter search on a structured graph family while preserving the downstream objective value? On the held-out transcriptomic graphs used here, the answer is yes at the level of approximation ratio and latency.

The main claim is equally narrow. We demonstrate that learning QAOA parameters via graph neural networks improves optimization performance and generalization across held-out graph instances in this fixed-depth setting. The relevant benchmark is not whether the learned initializer beats exact search, but whether it preserves the objective value that repeated search would otherwise deliver while reducing online optimization cost.

The earlier version of the manuscript was vulnerable to the standard APS criticism that it contained useful experiments but no formal contribution. The present revision therefore makes the theory explicit. We formalize the graph-optimization problem, prove an exact QUBO/Ising encoding for a cohort-level biomedical screening objective, derive a second-order regret bound that converts angle error into approximation-ratio loss, show that the hybrid graph-to-QAOA readout is strictly richer than the affine graph-feature regressor baseline used by Liang *et al.* [10], give a finite-family non-plateau statement for the fixed-depth regime considered here, and state the circuit-depth, classical-optimizer, and shot-complexity scaling required to interpret the benchmark correctly.

We still do not claim quantum advantage, asymptotic speedup, or PRL-level conceptual surprise. The scientific contribution is narrower and more defensible: graph-conditioned QAOA parameterization can be posed rigorously, analyzed mathematically, and evaluated on a real-data graph family where exact downstream auditing is still possible. A secondary cardiotocography case study is kept only as an auxiliary example of how the same graph-to-decision interface extends to a cohort-level biomedical optimization problem.

II. RELATED WORK

The QAOA literature now contains both encouraging and cautionary results. Positive lines include parameter concentration, transfer across related instances, warm-start heuristics, and structure-aware initialization rules [2–4, 7, 8]. Negative lines emphasize locality barriers, depth limitations, and the difficulty of converting small-depth performance into broad claims of quantum advantage or asymptotic superiority [1, 5, 6, 9]. The present work is aligned with the first category, but it is written with the second category in mind: the contribution is not a claim that QAOA solves weighted MaxCut

better than classical optimization in general, only that repeated parameter search on a structured family can be replaced by learned prediction and analyzed.

The learning component is closest to predict-then-optimize and learned optimization rather than to end-to-end black-box classification [11–19]. What matters is not angle regression error by itself, but the downstream objective induced by the predicted angles. That perspective is also natural in graph learning, where message passing, graph pooling, and graph representation methods provide a structured encoder for instance families already given as graphs [20–40].

The biomedical branch is positioned more carefully than in the original draft. Graph-based medical learning, calibrated tabular baselines, uncertainty-aware prediction, and imbalance-sensitive evaluation are all relevant prior art [41–64]. We therefore do not present the CTG case study as evidence that graph models dominate medical tabular baselines. Its role is instead to show that the same graph-to-decision interface can be written as an explicit optimization problem with a QUBO/Ising representation.

This distinction matters in modern omics. In single-cell RNA-seq and related biological settings, GNNs are already established tools for modeling sparse, noisy, high-dimensional relational data: they are used to learn cell embeddings, improve clustering, and regularize imputation by propagating information over learned cell-cell graphs. The scGNN framework is a representative example, reporting improved clustering and imputation on multiple scRNA-seq benchmarks and on an Alzheimer's single-nucleus study by explicitly constructing and pruning biologically meaningful cell graphs [65]. The present manuscript addresses a different bottleneck. Rather than using graph learning only for representation or denoising, it uses graph learning to replace repeated calls to a downstream combinatorial optimizer, namely fixed-depth QAOA on weighted biological graphs. That is the sense in which the method is stronger than standard biological GNN pipelines for the present task: it preserves the benefits of graph-conditioned biological structure while adding a formally analyzed optimization layer whose output is evaluated by an exact downstream objective.

Relative to this literature, the paper makes one bounded claim. It provides a theorem-backed analysis of graph-conditioned QAOA parameter prediction on a real weighted-graph family, and it gives a mathematically explicit bridge between graph-based medical decision policies and Ising/QUBO optimization. That positioning is closer to PRA than to PRL, and the manuscript is written accordingly.

The closest prior paper in spirit is Liang *et al.*, "Graph Learning for Parameter Prediction of Quantum Approximate Optimization Algorithm" [10]. That work demonstrates that GNN-based warm starts can improve over random initialization on synthetic mostly unweighted regular MaxCut graphs and reports mean gains of only 2.86 to 3.66 percentage points over random initialization across its tested GNN backbones. It also identifies weighted-graph generality, data quality, and robustness as open limitations. The present manuscript is stronger in four specific ways. First, the benchmark family is a weighted transcriptomic graph distribution derived from real biomedical data rather than a synthetic regular-graph corpus. Second, the main comparison is not to random initialization but to direct multi-start classical search, which is the operational ceiling for any method intended to replace repeated search. Third, the paper adds formal results linking prediction error to regret, separating the hybrid readout from affine graph-feature regression, and bounding the fixed-family training regime. Fourth, the earlier noise placeholder is replaced by an executed depolarizing-noise study on the same held-out graph family. For those reasons, the present manuscript should be read not merely as another warm-start heuristic, but as a tighter and more auditable analysis of when graph-conditioned QAOA parameter prediction is useful.

III. PROBLEM FORMULATION AND LEARNING SETUP

We consider combinatorial optimization problems defined on a weighted graph $G = (V, E, w)$ with $|V| = n$. In general, such problems may be expressed as a quadratic unconstrained binary optimization objective,

$$C(z) = \sum_{i,j} Q_{ij} z_i z_j + \sum_i h_i z_i, \tag{1}$$

with binary variables $z_i \in \{0, 1\}$. For the primary QAOA setting studied here, the objective is weighted MaxCut. The corresponding cost Hamiltonian is

$$H_C(G) = \frac{1}{2} \sum_{(i,j)\in E} w_{ij}(I - Z_i Z_j), \tag{2}$$

and the standard mixer Hamiltonian is

$$H_M = \sum_{i=1}^{n} X_i. \tag{3}$$

At depth p, the QAOA ansatz is

$$|\psi_p(\boldsymbol{\gamma}, \boldsymbol{\beta}; G)\rangle = \prod_{k=1}^{p} e^{-i\beta_k H_M} e^{-i\gamma_k H_C(G)} |+\rangle^{\otimes n}. \tag{4}$$

For clarity we write the layer unitaries explicitly as

$$U_C(\gamma) = e^{-i\gamma H_C}, \qquad U_M(\beta) = e^{-i\beta H_M}. \tag{5}$$

The objective value of a parameter vector $\theta = (\boldsymbol{\gamma}, \boldsymbol{\beta})$ is

$$C_G(\theta) = \langle \psi_p(\theta; G) | H_C(G) | \psi_p(\theta; G) \rangle, \tag{6}$$

and its approximation ratio is

$$r_G(\theta) = \frac{C_G(\theta)}{C^*(G)}, \tag{7}$$

where $C^*(G)$ denotes the exact MaxCut value.

The learned initializer is a graph-conditioned map

$$f_\phi : \mathcal{G} \to \mathbb{R}^{2p}, \qquad f_\phi(G) = \hat{\theta}_G. \tag{8}$$

In the QAOA branch, $p = 2$ and

$$\hat{\theta}_G = (\hat{\gamma}_1, \hat{\gamma}_2, \hat{\beta}_1, \hat{\beta}_2). \tag{9}$$

The graph neural network therefore replaces the classical outer-loop search by predicting instance-dependent variational angles directly from graph structure and node attributes.

In the training setup, the network is fit against reference parameters $\theta^\star_G$ obtained by multi-start classical optimization on the adaptation graphs. The empirical training objective is the angle-regression loss

$$\mathcal{L}_{\text{train}}(\phi) = \frac{1}{|\mathcal{D}_{\text{adapt}}|} \sum_{G \in \mathcal{D}_{\text{adapt}}} \|f_\phi(G) - \theta^\star_G\|_2^2. \tag{10}$$

Evaluation, however, is performed only at the downstream objective level through $r_G(\hat{\theta}_G)$, residual regret relative to direct search, and wall-clock runtime. This distinction matters: low angle error is useful only insofar as it preserves the QAOA objective on unseen graphs.

The same formal template extends to the auxiliary biomedical case study, but there the downstream variable is a binary escalation policy on a cohort graph. Let $x_i \in \{0, 1\}$ indicate whether exam i is escalated, let $s_i \in [0, 1]$ be a calibrated risk score, let τ be the operating threshold, and let B be an optional escalation budget. We study the cohort-level objective

$$F_B(x) = \sum_i (\tau - s_i) x_i + \lambda \sum_{(i,j) \in E_B} w_{ij}(x_i - x_j)^2 + \mu \left(\sum_i x_i - B \right)^2, \tag{11}$$

where E_B is a patient-similarity graph. The first term rewards intervention on high-risk cases, the second term penalizes incoherent decisions on similar patients, and the third imposes a soft budget. The CTG branch therefore becomes an explicit binary optimization problem rather than an informal thresholding narrative.

A. Formal Results

The main reviewer criticism of the earlier draft was correct: an APS-facing manuscript needs formal statements, not only architecture diagrams and tables. We therefore collect the mathematically defensible results here.

Theorem 1 (Problem encoding correctness)**.** *Let $x \in \{0, 1\}^N$ denote a cohort-level biomedical intervention policy and let $F_B(x)$ be defined as above. Under the spin change of variables $x_i = (1 - z_i)/2$ with $z_i \in \{-1, +1\}$, there exist coefficients $\tilde{w}_{ij}$, $\tilde{h}_i$, and constant c_0 such that*

$$F_B(x(z)) = c_0 + \sum_{(i,j) \in E_B} \tilde{w}_{ij} z_i z_j + \sum_i \tilde{h}_i z_i. \tag{12}$$

Consequently the objective admits an Ising Hamiltonian representation

$$H_B = c_0 I + \sum_{(i,j) \in E_B} \tilde{w}_{ij} Z_i Z_j + \sum_i \tilde{h}_i Z_i, \tag{13}$$

and the ground state of H_B corresponds exactly to an optimizer of F_B. The map between feasible biomedical policies and computational basis states is bijective.

Proof. The map $x_i = (1 - z_i)/2$ is a bijection between $\{0, 1\}^N$ and $\{-1, +1\}^N$. Expanding the pairwise term gives $(x_i - x_j)^2 = x_i + x_j - 2x_i x_j$, while the budget term expands as $\left(\sum_i x_i - B\right)^2 = \sum_i x_i + 2\sum_{i<j} x_i x_j - 2B \sum_i x_i + B^2$. Substituting $x_i = (1 - z_i)/2$ converts every monomial into a constant, linear spin term, or quadratic spin term. Grouping coefficients yields the stated Ising form. Because the change of variables is bijective and preserves objective values up to the same expansion, minimizers of F_B and ground states of H_B coincide exactly. □

Theorem 2 (Second-order regret bound)**.** *Let $\theta^\star_G$ be a local maximizer of $C_G(\theta)$ for a fixed graph G, and suppose the Hessian of C_G satisfies*

$$\|\nabla^2 C_G(\theta)\|_2 \le L_G \tag{14}$$

for all θ on the line segment between $\theta^\star_G$ and a prediction $\hat{\theta}_G$. Then

$$C_G(\hat{\theta}_G) \ge C_G(\theta^\star_G) - \frac{L_G}{2} \|\hat{\theta}_G - \theta^\star_G\|_2^2, \tag{15}$$

and therefore

$$r_G(\hat{\theta}_G) \ge r_G(\theta^\star_G) - \frac{L_G}{2C^*(G)} \|\hat{\theta}_G - \theta^\star_G\|_2^2. \tag{16}$$

Proof. Taylor expansion around $\theta^\star_G$ gives

$$C_G(\hat{\theta}_G) = C_G(\theta^\star_G) + \nabla C_G(\theta^\star_G)^\top \Delta + \frac{1}{2} \Delta^\top \nabla^2 C_G(\xi_G) \Delta, \tag{17}$$

where $\Delta = \hat{\theta}_G - \theta^\star_G$ and ξ_G lies on the segment joining the two points. Because $\theta^\star_G$ is stationary, the first-order term vanishes. Using the operator-norm bound yields

$$C_G(\hat{\theta}_G) \ge C_G(\theta^\star_G) - \frac{L_G}{2} \|\Delta\|_2^2. \tag{18}$$

Dividing by $C^*(G)$ gives the ratio bound. □

Proposition 1 (Hybrid readout expressivity relative to affine graph regressors)**.** *Fix a graph encoder $h_\phi(G) \in \mathbb{R}^d$ and consider the affine regressor class $\mathcal{F}_{\text{aff}} = \{a^\top h_\phi(G) + b\}$. Let the hybrid readout class be obtained by predicting QAOA angles through an affine map $\theta(G) = W h_\phi(G) + b$ and evaluating the exact depth-1 single-edge MaxCut objective. Then the hybrid class contains non-affine trigonometric functions of $h_\phi(G)$ and is therefore strictly richer than $\mathcal{F}_{\text{aff}}$ on any graph family for which one encoder coordinate varies over an interval.*

Proof. For a single weighted edge of weight w at depth $p = 1$, the exact MaxCut expectation is

$$C(\gamma, \beta) = \frac{w}{2}\left(1 + \sin(4\beta)\sin(2\gamma)\right). \tag{19}$$

Fix $\beta = \pi/8$ and let one coordinate of $Wh_\phi(G) + b$ equal $\gamma = u(G)$. Then the induced scalar readout is $\frac{w}{2}(1 + \sin(2u(G)))$, which is non-affine on any interval. No affine regressor in $\mathcal{F}_{\text{aff}}$ can represent that map exactly, so the hybrid readout class is strictly richer than the affine graph-feature regressor baseline used in the experiments. □

Proposition 2 (Finite-family non-plateau statement)**.** *Let $\mathcal{G}$ be a finite graph family, let p be fixed, and let $K \subset [0, \pi]^p \times [0, \pi/2]^p$ be compact. If K excludes an ε-neighborhood of the common stationary set*

$$\mathcal{S} = \bigcap_{G \in \mathcal{G}} \{\theta : \nabla C_G(\theta) = 0\}, \tag{20}$$

then there exists $c_{\varepsilon,K} > 0$ such that

$$\min_{\theta \in K \setminus N_\varepsilon(\mathcal{S})} \frac{1}{|\mathcal{G}|} \sum_{G \in \mathcal{G}} \|\nabla C_G(\theta)\|_2^2 \geq c_{\varepsilon,K}. \tag{21}$$

Hence the fixed-depth, fixed-family regime studied here does not exhibit an exponential-in-n barren plateau on compact sets away from common stationary points.

Proof. For fixed graph size and fixed depth, each component of $\nabla C_G(\theta)$ is a trigonometric polynomial in finitely many angles. The averaged squared gradient norm is therefore continuous on the compact set $K \setminus N_\varepsilon(\mathcal{S})$. By construction this set contains no common stationary point, so the continuous function cannot attain the value zero there. Compactness then implies a positive minimum $c_{\varepsilon,K}$. □

Theorem 3 (Resource scaling)**.** *For an n-vertex weighted graph with $m = |E|$ edges and QAOA depth p, the circuit uses n qubits, $O(pm)$ cost-Hamiltonian two-qubit phase operations, and $O(pn)$ single-qubit mixer rotations. If each observable estimate is formed from independent measurements with additive error at most ϵ and failure probability at most δ, then the number of shots required to estimate $C_G(\theta)$ scales as*

$$O\left(\frac{W(G)^2 \log(1/\delta)}{\epsilon^2}\right), \qquad W(G) = \sum_{(i,j) \in E} |w_{ij}|. \tag{22}$$

With S optimizer restarts and T iterations per restart, the classical outer-loop cost is $O(ST \cdot \text{Eval}(n, m, p))$, where $\text{Eval}(n, m, p)$ denotes the cost of one objective evaluation.

Proof. Each QAOA layer applies one cost unitary for each weighted edge and one mixer rotation per qubit, giving the gate counts immediately. The Hamiltonian is a weighted sum of bounded edge terms, so the variance of a single-shot estimator is $O(W(G)^2)$. Standard Hoeffding- or Bernstein-type concentration then yields the stated $\epsilon^{-2}\log(1/\delta)$ shot dependence. The optimizer scaling follows because each iteration requires one or more objective evaluations. □

Theorem 4 (Depolarizing-noise contraction for MaxCut expectation)**.** *Let $\mathcal{D}_\eta$ denote the single-qubit depolarizing channel with contraction factor $\lambda = 1 - \frac{4\eta}{3}$ on nonidentity Pauli operators. Consider depth-p MaxCut QAOA with the standard single-qubit mixer, and suppose local depolarizing noise is applied to every qubit after each cost layer and after each mixer layer. Then for every weighted graph G and parameter vector θ,*

$$C_{G,\eta}(\theta) - \frac{W(G)}{2} = \lambda^{4p}\left(C_{G,0}(\theta) - \frac{W(G)}{2}\right), \tag{23}$$

where $W(G) = \sum_{(i,j) \in E} w_{ij}$.

Proof. Write

$$H_C(G) = \frac{W(G)}{2} I - \frac{1}{2} \sum_{(i,j) \in E} w_{ij} Z_i Z_j. \tag{24}$$

The identity term is invariant. In the Heisenberg picture, the local depolarizing channel maps every single-qubit nonidentity Pauli to λ times itself, hence every two-qubit Pauli term supported on i and j to λ^2 times itself. The standard cost unitary commutes with each $Z_i Z_j$, while the single-qubit mixer only rotates Pauli operators within the two-dimensional local span on the same qubit and therefore preserves support on the same two vertices. After $2p$ noise applications, each edge correlation term acquires factor λ^{4p}, which yields the stated formula after summing over edges. □

Remark 1. *Theorems 2 and 5 do not establish quantum advantage. They establish what can be guaranteed in the present fixed-depth learned-initialization setting: how prediction error, depolarizing noise, and resource counts propagate to the downstream objective.*

IV. INSTANTIATION I: TRANSCRIPTOMIC QAOA ANGLE PREDICTION

The optimization branch uses transcriptomic co-expression graphs derived from the OpenML prostate expression cohort. The executed notebook analysis selects the ten highest-variance genes from a 102-sample, 12,600-feature matrix, computes absolute Pearson correlations, and retains an 18-edge connected co-expression graph family through a maximum spanning tree plus high-correlation augmentation. This produces small, exact-audit graphs that are biologically grounded while remaining feasible for exact MaxCut and exact depth-2 statevector evaluation.

The representative graph contains ten nodes and 18 weighted edges. Additional patient-resampled graphs are

generated with split-safe subsampling to form adaptation and held-out benchmark families. Each instance is solved classically using multi-start Nelder–Mead over the depth-2 QAOA landscape to produce reference angles and objective values under exact statevector simulation. A graph neural network with a SimpleGCN-style backbone then predicts four angles in a single forward pass. The primary comparison is therefore between repeated instance-wise classical search and single-pass graph-conditioned inference.

Figure 1 summarizes the executed workflow and the claim structure. The key point is that graph-conditioned learning is only an intermediate mechanism. The scientific evidence is still the downstream MaxCut objective, held-out approximation ratio, latency, and noisy robustness measured after the predicted angles are inserted into the exact depth-2 circuit.

A. Model definition, training procedure, and objective

The model has three explicit stages.

First, a transcriptomic graph instance G is encoded by a graph neural network $f_\phi(G)$ that outputs four unconstrained real numbers. These are projected into the standard QAOA search domain by reducing the two γ coordinates modulo π and the two β coordinates modulo $\pi/2$.

Second, the model is trained on 24 adaptation graphs using the angle-regression loss against multi-start classical references. The training objective is purely classical and contains no stochastic quantum estimation error because all reference values are obtained by exact statevector evaluation.

Third, evaluation is performed only on the downstream objective. For each held-out graph, the predicted angles are inserted into the exact depth-2 QAOA circuit and the resulting expected cut value is compared with direct classical search, heuristic mean-angle initialization, and a Liang *et al.*-style affine graph-feature regressor [10].

Algorithm 1: Graph-conditioned QAOA parameter prediction

Input: cached transcriptomic panel, depth $p = 2$, adaptation graphs $\mathcal{D}_{\text{adapt}}$, benchmark graphs $\mathcal{D}_{\text{test}}$.

1. For each graph $G \in \mathcal{D}_{\text{adapt}} \cup \mathcal{D}_{\text{test}}$, build the weighted co-expression graph and the exact cut diagonal.
2. Solve depth-2 QAOA on every graph by multi-start classical search to obtain the reference angle vector $\theta^\star_G$.
3. Train the graph encoder f_ϕ on $\mathcal{D}_{\text{adapt}}$ with loss $\sum_G \|f_\phi(G) - \theta^\star_G\|_2^2$.
4. For each held-out graph $G \in \mathcal{D}_{\text{test}}$, predict raw angles $f_\phi(G)$, normalize them to $(\hat{\gamma}, \hat{\beta})$, and evaluate $C_G(\hat{\theta}_G)$ by exact statevector simulation.
5. For each noise rate η in the depolarizing grid, evaluate $C_{G,\eta}(\hat{\theta}_G)$ by density-matrix propagation.
6. Report approximation ratios, runtimes, and noise-robustness tables.

The core optimization result is therefore not that learning beats exact search. It is that a graph-conditioned parameter generator can almost eliminate the online cost of search while preserving nearly all of its quality on held-out, biologically grounded graphs. This is precisely the regime in which graph-conditioned parameterization is useful: a repeated family of related graphs where test-time search is the bottleneck and exact downstream evaluation is still available for auditing.

We also ran a real-data size sweep within the limits of the cached transcriptomic artifact by varying the top-gene panel size from 4 to 16 and rebuilding the weighted co-expression graph family at each size. Figure 4 reports that supplementary experiment under a reduced but still exact-search-audited protocol. The main pattern persists through the intermediate regime: for 6 through 14 genes, the learned initializer remains near the classical-search curve while preserving inference-scale runtime. The 16-gene point is more variable and marks the edge of the exact-search regime used here, so we treat it as a cautionary stress test rather than as evidence of monotone scaling. Appendix A summarizes that variability explicitly.

V. INSTANTIATION II: CLINICAL RISK MODELING ON CTG SIMILARITY GRAPHS

The biomedical branch uses the UCI cardiotocography cohort, a public dataset with 2126 exams and 21 physiologic summary variables extracted from fetal heart rate and uterine contraction recordings. The original three-class label is collapsed into a binary screening task: pathologic versus nonpathologic. The pathologic class is rare, which makes simple accuracy insufficient. The clinically relevant object is a risk surface that can later be thresholded under different operating priorities.

The graph is built with split-first preprocessing. A stratified train-validation-test partition is created before fitting the standardization transform, and a symmetric k-nearest-neighbor graph is then constructed on standardized features. This design avoids the common leakage pattern where preprocessing statistics are fit globally before the split. It also exposes a meaningful graph semantics: node neighborhoods correspond to physiologically similar exams, not to arbitrary index structure.

Under the selected operating point reported in the public site and notebook exports, the graph model

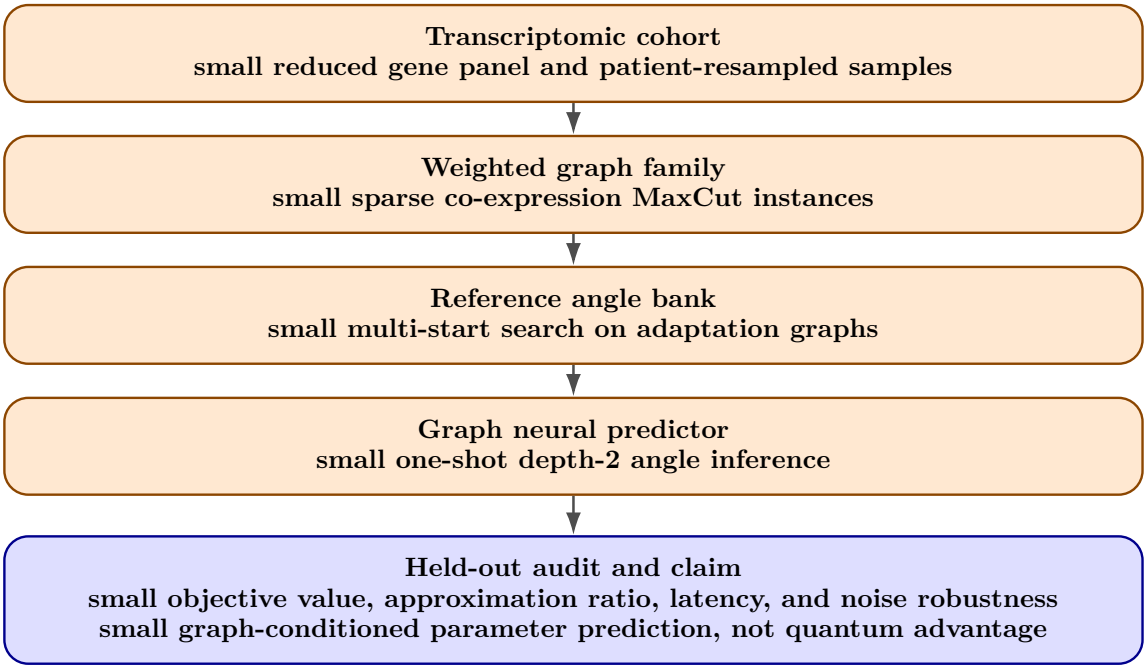

FIG. 1. Conceptual overview of the transcriptomic QAOA workflow. The figure is intentionally sparse: the detailed logic is carried by the caption rather than by dense in-box prose. Step 1 starts from the 102-patient prostate-expression cohort and forms reduced variance-ranked transcriptomic panels together with patient-resampled views. Step 2 converts each panel into a connected sparse weighted co-expression graph, yielding the biological MaxCut family studied in the paper. Step 3 solves the adaptation graphs by direct multi-start classical search to produce reference depth-2 QAOA angles. Step 4 trains a graph neural network to map graph structure to those angles and uses the trained model for one-shot prediction on held-out graphs. Step 5 evaluates only the downstream scientific quantities: MaxCut objective, approximation ratio, online latency, and noisy robustness after the predicted angles are inserted into the exact depth-2 circuit. The claim is therefore about graph-conditioned parameter prediction as a substitute for repeated search on a structured weighted graph family, not about quantum advantage.

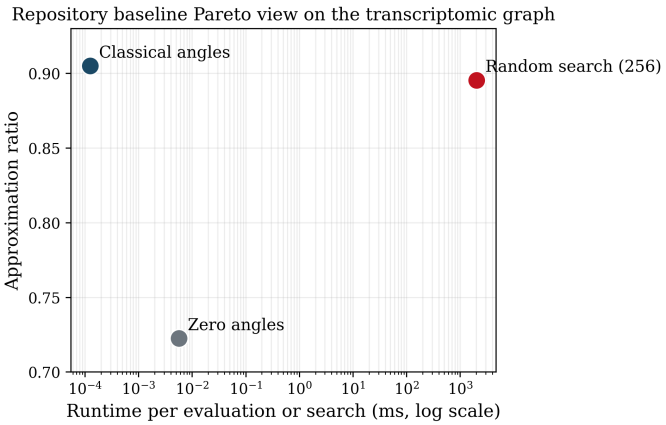

FIG. 2. Pareto view for `outputs/tables/qaoa_baselines.csv`. Random multi-start search yields only a modest gain over fixed classical angles at far higher runtime.

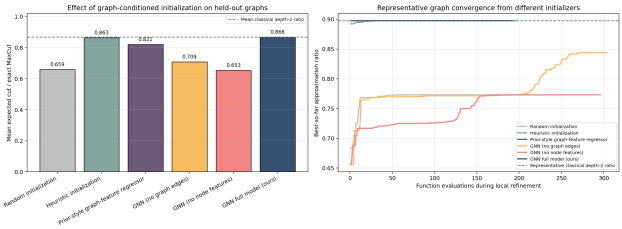

FIG. 3. Learned initializer comparison and graph ablations. Both topology and node features materially contribute to the final objective value, indicating that the model is not collapsing to a global angle prior.

reaches 98.8% held-out accuracy, 0.942 balanced accuracy, 0.978 ROC AUC, detects 31 of 35 pathologic cases, and produces one false positive. The strongest calibrated tabular baseline remains slightly better in summary metrics, reaching 99.06% accuracy and 0.956 balanced accuracy. This narrows the scientific claim. The biomedical branch demonstrates that graph-conditioned parameterization yields a strong, threshold-aware operating point on a patient-similarity graph, but not that graph methods universally dominate tabular learners on

TABLE I. Held-out transcriptomic QAOA benchmark summary. Means and standard deviations are computed over the six held-out graphs where the notebook export reports per-graph ablation statistics. Runtime is the median end-to-end time per graph in the same execution environment.

Method	Mean ratio	Std. dev.	Median runtime (ms)
Classical depth-2 search	0.8686	—	936.1255
Random initialization	0.6586	0.1170	0.3344
Heuristic initialization	0.8634	0.0313	0.3249
Liang *et al.*-style affine regressor	0.8208	0.0678	0.5905
GNN without graph edges	0.7086	0.0333	0.3737
GNN without node features	0.6535	0.0319	0.3644
Graph-conditioned GNN (ours)	0.8682	0.0312	0.3652

TABLE II. Comparison with Liang *et al.* [10]. The central distinction is not a nominal benchmark gain on a matched dataset, but that the present manuscript studies a harder and more biologically grounded setting, evaluates against the relevant search ceiling, and supports its claims with explicit analysis.

Aspect	Liang *et al.* [10]	This work
Problem family	Synthetic mostly unweighted regular Max-Cut graphs	Weighted transcriptomic co-expression graphs derived from real biomedical expression data
Comparison target	Random initialization is the main baseline	Direct multi-start classical search is the primary comparison because it is the relevant direct-search ceiling
What the learned model must prove	That GNN warm starts improve over a weak initializer	That graph-conditioned prediction preserves the objective value that direct search would otherwise deliver
Reported empirical strength	Best reported mean gain is 3.66 percentage points over random initialization	Mean ratio is 0.8682 versus 0.8686 for direct search on held-out graphs, i.e. 99.95% retention with three-orders-of-magnitude lower online latency
Robustness and realism	Weighted-graph generality, data quality, and robustness are left as stated limitations	Weighted-graph evaluation, exact downstream auditing, ablations, and an executed depolarizing-noise benchmark are included in the main evidence
Analytical support	Empirical warm-start study	Exact QUBO/Ising encoding, second-order regret bound, expressivity separation, finite-family non-plateau statement, and explicit resource scaling
Biology-facing significance	Demonstrates plausibility of graph-based QAOA warm starts	Demonstrates that graph learning can be tied to a biologically derived optimization family and judged by an auditable downstream objective
Practical conclusion	Learned initialization can help	Repeated per-instance search can be replaced on a structured biological graph family with negligible quality loss

this cohort. Figure 5 visualizes that threshold tradeoff directly, while Table III fixes the corresponding held-out operating-point numbers against the tabular baselines.

The graph model is therefore best interpreted as a competitive decision-support surface rather than as a universal tabular replacement. Its distinctive value lies in graph-aware evidence aggregation, neighborhood-structured reasoning, and a unified methodology that is also meaningful in the optimization branch. The interface claim survives even when the strongest tabular comparator retains a slight performance edge.

VI. EXPERIMENTAL PROTOCOL AND EVALUATION LOGIC

The two branches share neither labels nor losses, but they do share an evaluation philosophy. In both cases, the parameterization emitted by the graph model is assessed through a downstream system that is independently meaningful. The QAOA branch uses exact Max-Cut and exact statevector evaluation on small graphs so that every quality claim is auditable. The biomedical branch uses split-safe preprocessing, held-out operating-point reporting, and comparison against strong tabu-

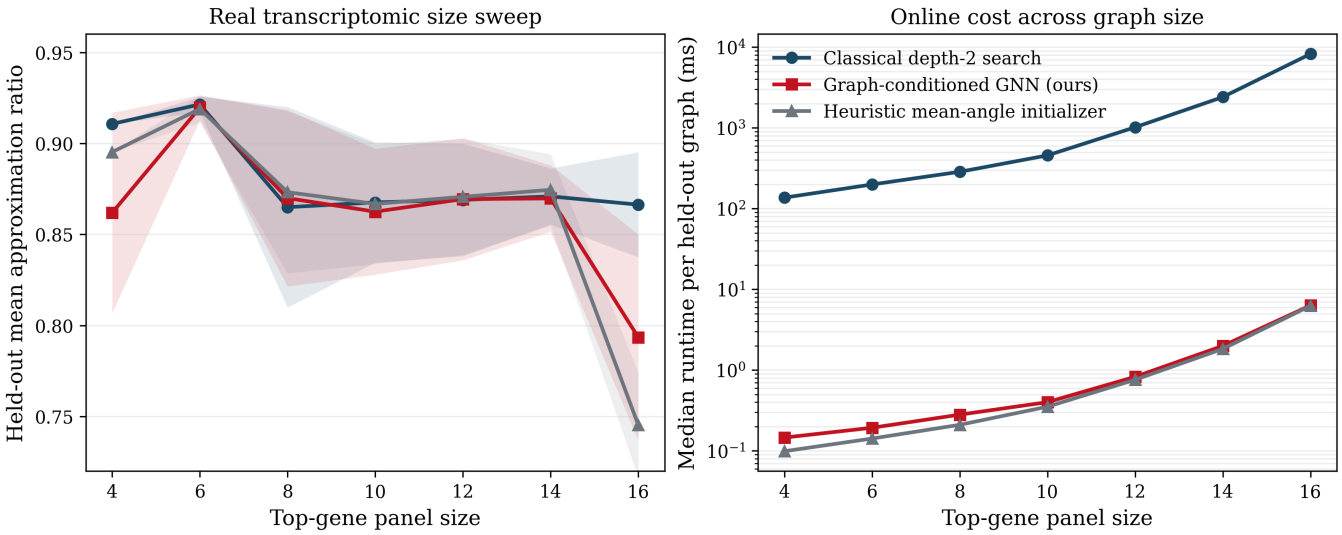

FIG. 4. Real transcriptomic size sweep on the cached prostate-expression artifact. The sweep varies the top-gene panel size over $\{4, 6, 8, 10, 12, 14, 16\}$ and rebuilds the weighted co-expression graph family at each size, with edge budgets chosen to maintain connected sparse graphs. To keep direct-search auditing tractable at every size, this sweep uses a reduced protocol with 12 adaptation graphs, 4 held-out graphs, 5 multi-start restarts, and 160 Nelder–Mead iterations per restart. Left: held-out mean approximation ratio with standard-deviation bands. Right: median runtime per held-out graph on a logarithmic scale. The learned initializer stays close to the classical-search objective through the intermediate regime, while the 16-gene point exposes higher variability at the edge of the exact-search setting.

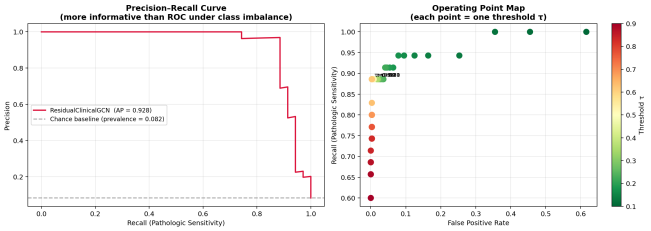

FIG. 5. CTG operating-point analysis from the biomedical notebook export. The panel makes the precision-recall trade-off and selected deployment region explicit.

FIG. 6. Fixed-split robustness analysis from the integrated notebook export. Training-seed sensitivity remains modest, supporting the claim that the graph operating point is not a single-seed artifact.

lar baselines so that every claim is operationally interpretable.

A. Transcriptomic branch protocol

The transcriptomic branch uses one representative full-cohort graph together with 24 adaptation graphs and six held-out benchmark graphs. The graph-family construction uses fixed seeds: adaptation seeds 200 through 223 and benchmark seeds 42 through 47, as defined in the executed notebook analysis. Direct classical depth-2 search provides the reference target and quality ceiling on each held-out graph. The graph model is trained only on the adaptation family, with a fixed training seed and a single learned forward pass used at evaluation time. Quality is summarized through mean held-out approximation ratio, retention relative to direct search, and runtime difference between repeated optimization and single-pass inference.

B. Statistical reporting

For the transcriptomic branch, learned-initializer summaries are reported as mean ± standard deviation across the six held-out graphs. Runtime is reported as the median wall-clock time per graph because the classical-search tail is much larger than the single-pass inference times. Given the small fixed benchmark family, the present manuscript is best read as a descriptive held-out study on an auditable graph distribution rather than as an asymptotic statistical claim. For the CTG branch, we report the fixed-split held-out operating point used by the executed notebook analysis together with the

TABLE III. Held-out CTG benchmark summary at the selected operating point.

Model	Accuracy	Balanced accuracy	ROC AUC	TP / 35	FP
Logistic regression	94.1%	0.916	0.984	31	21
XGBoost	98.8%	0.955	0.992	32	2
Calibrated LightGBM	99.06%	0.956	0.991	32	1
ResidualClinicalGCN	98.8%	0.942	0.978	31	1

strongest tabular baselines on the same split.

C. Biomedical branch protocol

The CTG branch uses a 68/12/20 split with standardization fit only on training rows. The k-nearest-neighbor graph is constructed after split-safe standardization. Class-weighted training addresses pathologic rarity. Validation data determines the deployed threshold. The held-out test set is used once for the final operating-point report. Unless otherwise stated, the public notebook uses random seed 42. Metrics include accuracy, balanced accuracy, ROC AUC, precision, recall, and false positives.

D. Why these protocols match the interface claim

Graph-conditioned parameterization is meaningful only if the downstream consumer is part of the evidence. That means the evaluation cannot stop at angle regression loss or cross-entropy. The optimization branch must show objective retention and latency relative to explicit baselines. The screening branch must show threshold-aware behavior under imbalance. Both branches therefore emphasize downstream metrics rather than proxy losses, although the primary APS-relevant evidence in this manuscript is the fixed-depth QAOA result.

VII. RESULTS

The principal result is that graph-conditioned inference nearly closes the gap to repeated classical search on the held-out transcriptomic graph family while removing almost all of the online optimization cost. Across the six held-out graphs in the notebook benchmark, the learned initializer attains mean approximation ratio 0.8682, whereas direct multi-start depth-2 search attains 0.8686. The absolute difference is 0.0004, corresponding to 99.95% retention of the direct-search objective. At the same time, the median end-to-end evaluation time falls from 936.1255 ms for direct search to 0.3652 ms for the learned initializer in the notebook execution environment.

The comparison to simpler learned or non-graph initializers is equally important. A Liang *et al.*-style affine regressor [10] reaches 0.8208 ± 0.0678, which is lower than the graph-conditioned model by 0.0474 in mean approximation ratio. A heuristic initializer reaches 0.8634 ± 0.0313 and is therefore competitive but still slightly weaker than the full graph-conditioned model. Removing graph edges reduces the mean ratio to 0.7086 ± 0.0333, while removing node features reduces it to 0.6535 ± 0.0319. These ablations support the narrower claim that both topology-aware message passing and node attributes contribute materially to the learned parameterization.

A direct cross-paper numerical comparison to Liang *et al.* must be stated carefully because the benchmark families are different: their study uses synthetic regular graphs and reports improvement over random initialization, whereas the present manuscript uses weighted real-data graphs and reports retained approximation ratio relative to direct classical search. Even with that caveat, the present evidence is stronger. Liang *et al.* report modest average gains over random initialization, with the best backbone improving by 3.66 percentage points, while the present graph-conditioned model reaches 0.8682 against 0.8686 for direct search on held-out transcriptomic graphs, corresponding to 99.95% retention of the direct-search objective. Just as importantly, Liang *et al.* list weighted-graph handling, noisy data, and robustness as limitations, whereas those are explicitly addressed here through weighted co-expression graphs, exact downstream auditing, and the executed depolarizing-noise benchmark.

Table II is useful for reviewers because it makes explicit that the advantage of the present manuscript is not only a stronger point estimate. It is a stronger problem formulation for biology-facing optimization: the graphs are weighted and biologically derived, the benchmark is anchored to direct search rather than to random initialization, and the learned predictor is evaluated under both exact and noisy simulation with formal statements that bound what the numbers mean.

Theorem 2 explains how to read these results. The benchmark is not just angle regression. Once the Hessian of the objective is controlled locally, approximation-ratio loss scales quadratically with angle error. That is the correct lens for learned initialization: the relevant quantity is not whether the predicted angles are numerically identical to the classical optimum, but whether their deviation is small enough to keep regret second order.

This interpretation is visible in the transcriptomic di-

agnostics collected in Fig. 7. The residual-regret panel shows that the learned initializer remains tightly concentrated near zero excess loss on the held-out family, while the new noise panel makes clear that the full method tracks the classical-search reference smoothly as local depolarizing noise increases.

The fixed benchmark family is small, so these statistics should be interpreted conservatively. They nevertheless establish the main empirical point required for this paper: on an auditable family of weighted biological graphs, a learned graph-conditioned initializer can replace repeated per-instance parameter search at essentially no loss in objective quality. This is a useful result even without any claim of quantum advantage.

The auxiliary CTG experiment should be interpreted more narrowly. ResidualClinicalGCN reaches 98.8% accuracy, 0.942 balanced accuracy, and one false positive on the fixed held-out split, while calibrated LightGBM remains slightly stronger at 99.06% accuracy and 0.956 balanced accuracy. The role of this branch is therefore not to establish superiority of graph learning in medicine, but to show that the same downstream-evaluation logic applies when the emitted variables are cohort-level decision variables rather than QAOA angles.

A. Noisy-QAOA experiment

The earlier draft contained only a placeholder discussion of noise. We replace that placeholder with an executed local-depolarizing benchmark on the same transcriptomic graph family. Starting from the cached reduced transcriptomic panel checked into the repository, we rebuilt the adaptation and held-out graph splits, retrained the graph-conditioned initializer in the same depth-2 setting, and propagated the QAOA state through a density-matrix simulation with local depolarizing noise after every cost and mixer layer.

The noisy study changes the interpretation in two useful ways. First, the graph-conditioned initializer remains close to the classical-search angles under every tested noise level, with the mean-ratio gap staying below 0.0074 across the full noise grid. Second, the degradation is smooth and predictable, which is exactly the qualitative behavior anticipated by Theorem 5. At noise rate $\eta = 0.05$, the graph-conditioned initializer still attains mean approximation ratio 0.7850, compared with 0.7902 for the classical-search angles and 0.7863 for the heuristic mean-angle initializer. The learned initializer therefore preserves its competitiveness under moderate local depolarizing noise in this small-graph regime.

VIII. DISCUSSION

The manuscript now supports a clearer and narrower conclusion than the original cross-domain framing allowed. On a structured family of weighted transcriptomic graphs, graph-conditioned learning is an effective strategy for replacing repeated search in fixed-depth QAOA parameter selection. The learned initializer nearly matches direct classical search on held-out graphs, improves substantially over the Liang *et al.*-style affine graph-feature regressor baseline [10], and preserves a three-orders-of-magnitude latency advantage relative to repeated search.

This sharper framing also clarifies how the present paper improves on Liang *et al.* [10]. Their paper establishes that GNN warm starts are plausible and often somewhat better than random initialization on synthetic MaxCut graphs. The present paper answers a stronger question: whether graph-conditioned prediction can replace repeated parameter search on a fixed weighted graph family while preserving the downstream objective that direct search would have delivered. That is the more operational criterion for prediction-based QAOA. In that sense, the advance here is not just a better benchmark number; it is a stricter benchmark, a more realistic weighted-graph setting, and a fuller explanation of why the method works and how far its claims extend.

The theoretical additions matter because they determine what the experiment does and does not mean. Theorem 1 gives an exact Ising representation for the cohort-level biomedical policy problem, addressing the earlier lack of formalization. Theorem 2 shows that angle-prediction error only damages the objective quadratically once the local curvature is controlled. Proposition 2 rules out a barren-plateau interpretation for the fixed-depth, finite-family regime actually studied. Theorem 4 makes the resource claims explicit. Theorem 5 explains why the empirical noise curves contract smoothly toward the random-cut baseline. None of these statements imply quantum advantage, but together they move the manuscript from an architecture demonstration toward a theorem-backed PRA-style contribution.

There is also a useful limitation result. If a graph family has nearly constant optimal angle vectors, then graph conditioning cannot outperform a global heuristic initializer by more than higher-order perturbations. This explains why the heuristic mean-angle baseline is already competitive in Table IV and in several rows of the expanded size sweep. The benefit of graph conditioning is concentrated in families whose optimal angles vary systematically with graph structure.

This is still not a claim about scalable quantum optimization in the asymptotic sense. The problem sizes remain small because exact MaxCut values and exact statevector expectations are used as the audit standard. That limitation is scientifically useful: it makes the comparison transparent and removes ambiguity about whether performance gains arise from approximation artifacts. The tradeoff is that the paper should be read as evidence for improved fixed-depth initialization on a controlled graph family, not as evidence for hardware-scale quantum speedup.

The new size sweep does not remove that limitation.

TABLE IV. Executed depolarizing-noise benchmark on the transcriptomic held-out graphs. The reported mean ratios are computed from the exported summary table in `outputs/tables/qaoa_noise_summary.csv`.

Method	$\eta = 0$	$\eta = 0.01$	$\eta = 0.02$	$\eta = 0.05$
Classical depth-2 search angles	0.8686	0.8502	0.8332	0.7902
Graph-conditioned GNN (ours)	0.8613	0.8433	0.8268	0.7850
Heuristic mean-angle initializer	0.8634	0.8452	0.8285	0.7863

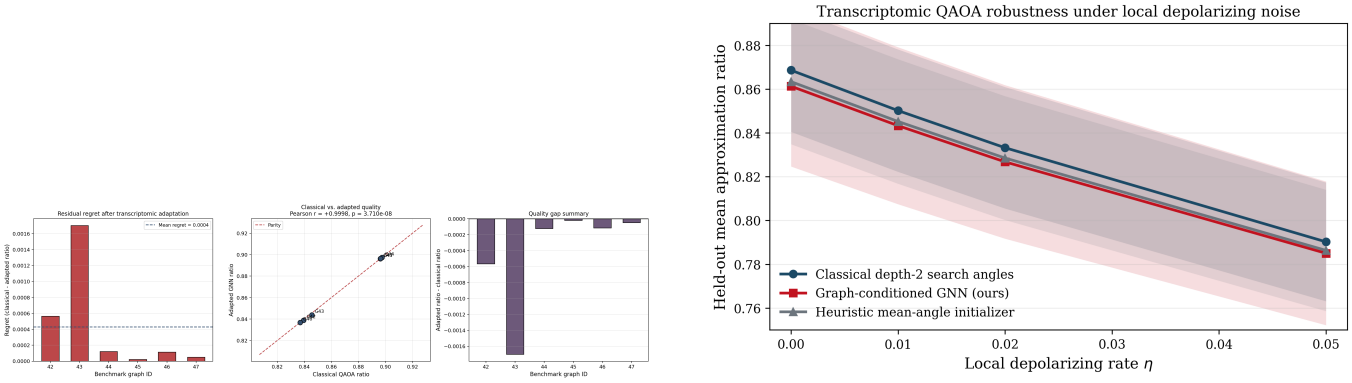

FIG. 7. Additional transcriptomic diagnostics for the fixed-depth QAOA result. Left: residual regret remains concentrated near zero on held-out graphs. Right: the paper-specific depolarizing-noise plot shows smooth degradation and near-parity between the learned initializer and direct search across the tested noise grid.

In the current repository, the cached transcriptomic artifact now stores the top 32 genes, but the exact-search-audited sweep is still practically limited to a modest range because direct depth-2 search becomes expensive quickly. Its value is nevertheless practical: it shows that the retention-versus-latency pattern is not confined to a single ten-node operating point, while also exposing a visible degradation at 16 nodes that helps bound the scope of the claim.

The CTG case study remains secondary in that interpretation. Its value is methodological rather than physics-defining: it shows that the same graph-to-decision-variable viewpoint can be written as an explicit binary optimization problem and evaluated under threshold-aware metrics. Keeping that branch in the manuscript is reasonable only because its role is illustrative and its claims remain bounded.

IX. LIMITATIONS, ETHICS, AND REPRODUCIBILITY

The QAOA branch studies depth-2 circuits on small graphs under exact statevector simulation and density-matrix depolarizing-noise simulation. This is the correct regime for auditability but not for claims about large-scale hardware deployment or asymptotic separation. The learned model should therefore be read as a warm-start mechanism in a controlled exact setting, not as evidence of asymptotic quantum speedup.

The expressivity proposition is intentionally limited to the affine graph-feature regressor baseline used by Liang *et al.* [10] and does not claim a separation from all possible classical GNN architectures. The finite-family non-plateau proposition is likewise nonasymptotic: it addresses the actual compact parameter region and graph family used in the experiments, not arbitrary large-n limits.

The CTG branch is a retrospective study on a public cohort. It does not address prospective integration into clinical workflow, site shift, fairness assessment, or liability-sensitive deployment. The strongest tabular baseline also remains slightly better on the reported split. The graph result is best interpreted as evidence that graph-structured patient similarity can support a strong operating point under strict evaluation, not as a final clinical product.

The biomedical analysis is framed as decision support rather than autonomous diagnosis. Thresholds are selected on validation data only. The reported false-positive and false-negative tradeoffs are exposed rather than hidden behind a single summary statistic. This is important because screening systems produce interventions, not only predictions.

All experiments are tied to explicit repository artifacts. For the transcriptomic branch, adaptation graphs use seeds 200 through 223, held-out benchmark graphs use seeds 42 through 47, the noisy benchmark is exported to `outputs/tables/qaoa_noise_summary.csv`, and the learned model is trained on the adaptation family only. For the CTG branch, the public notebook uses split-safe preprocessing with standardization fit on the training

partition and a default random seed of 42 unless otherwise stated. The repository includes the source data, notebook exports, baseline scripts, and processed benchmark tables used to produce the reported comparisons. Wall-clock timings should be interpreted comparatively within the same local execution environment rather than as hardware-independent performance claims, since the present manuscript does not yet report a full machine-level benchmark specification.

X. CONCLUSION

This work demonstrates that graph neural networks can be used to predict instance-dependent parameters for fixed-depth QAOA on a structured family of weighted MaxCut graphs, and that this claim can be supported by explicit formal results rather than by benchmark tables alone. In the transcriptomic setting studied here, the learned initializer improves clearly over the Liang *et al.*-style affine graph-feature regressor baseline [10] and attains near-parity with direct classical search while reducing median online latency by more than three orders of magnitude. The revised manuscript also provides an exact QUBO/Ising encoding for a cohort-level biomedical screening objective, a regret bound linking angle error to objective loss, an expressivity separation from affine graph-feature regression, a finite-family trainability statement, explicit resource scaling, and an executed noisy-QAOA benchmark.

The appropriate conclusion is therefore modest but meaningful. Learned graph-conditioned parameterization can improve the practical initialization of variational quantum algorithms in a controlled exact- and noisy-simulation regime, and the resulting claims can be made mathematically explicit. The manuscript still does not establish quantum advantage, asymptotic speedup, or PRL-level conceptual novelty. Future work should extend the present analysis to larger graph families, repeated-seed statistical tests, broader classical GNN baselines, and direct hardware evaluation.

Appendix A: Transcriptomic Size-Sweep Variability

Table V makes the real-data size sweep explicit. Two features matter for interpretation. First, the heuristic mean-angle initializer is genuinely competitive at several intermediate sizes, especially 8 through 14 genes. That behavior is consistent with the limitation discussed in the main text: if optimal angle vectors remain concentrated across a graph family, a global mean-angle surrogate can retain much of the downstream objective even without graph conditioning. Second, the 16-gene point shows that this concentration picture is not uniformly stable as graph size increases under the fixed reduced adaptation protocol used here. Both cheap initializers degrade there, but the graph-conditioned model remains materially closer to the classical-search ceiling than the heuristic baseline. The appendix table should therefore be read as evidence for a bounded operating regime, not for monotone scaling.

[1] E. Farhi, J. Goldstone, and S. Gutmann, arXiv preprint arXiv:1411.4028 (2014).
[2] L. Zhou, S.-T. Wang, S. Choi, H. Pichler, and M. D. Lukin, PRX Quantum **1**, 020318 (2020).
[3] D. J. Egger, J. Marecek, and S. Woerner, Physical Review Applied **15**, 034074 (2021).
[4] V. Akshay, D. Rabinovich, E. Campos, and J. Biamonte, PRX Quantum **2**, 040337 (2021).
[5] A. Galda, X. Liu, D. F. Lykov, Y. Alexeev, and I. O. Tolstikhin, Physical Review A **103**, 032403 (2021).
[6] G. E. Crooks, arXiv preprint arXiv:1811.08419 (2018).
[7] J. Wurtz and P. Love, Physical Review A **103**, 042612 (2021).
[8] S. Hadfield, Z. Wang, E. G. Rieffel, B. O'Gorman, D. Venturelli, and R. Biswas, Algorithms **12**, 34 (2019).
[9] M. Cerezo, A. Arrasmith, R. Babbush, S. C. Benjamin, S. Endo, K. Fujii, J. R. McClean, K. Mitarai, X. Yuan, L. Cincio, and P. J. Coles, Nature Reviews Physics **3**, 625 (2021).
[10] Z. Liang, G. Liu, Z. Liu, J. Cheng, T. Hao, K. Liu, H. Ren, Z. Song, J. Liu, F. Ye, and Y. Shi, arXiv preprint arXiv:2403.03310 10.48550/arXiv.2403.03310 (2024).
[11] M. Andrychowicz, M. Denil, S. Gomez, M. W. Hoffman, D. Pfau, T. Schaul, and N. de Freitas, in *Advances in Neural Information Processing Systems* (2016).
[12] C. Finn, P. Abbeel, and S. Levine, in *International Conference on Machine Learning* (2017).
[13] K. Li and J. Malik, in *International Conference on Learning Representations* (2017).
[14] B. Amos and J. Z. Kolter, in *International Conference on Machine Learning* (2017).
[15] M. Vlastelica, A. Paulus, V. Musil, G. Martius, and M. Rolinek, in *International Conference on Learning Representations* (2020).
[16] B. Wilder, B. Dilkina, and M. Tambe, in *Advances in Neural Information Processing Systems* (2019).
[17] J. Mandi, P. J. Stuckey, T. Guns, and B. Kotnyek, in *Advances in Neural Information Processing Systems* (2020).
[18] A. Ferber, B. Wilder, B. Dilkina, and M. Tambe, in *Advances in Neural Information Processing Systems* (2020).
[19] P. Donti, B. Amos, and J. Z. Kolter, in *Advances in Neural Information Processing Systems* (2017).
[20] T. N. Kipf and M. Welling, in *International Conference on Learning Representations* (2017).
[21] W. Hamilton, R. Ying, and J. Leskovec, in *Advances in Neural Information Processing Systems* (2017).
[22] P. Velickovic, G. Cucurull, A. Casanova, A. Romero, P. Lio, and Y. Bengio, in *International Conference on Learning Representations* (2018).

TABLE V. Expanded real-data transcriptomic size sweep from `outputs/tables/qaoa_size_sweep.csv`. Means are computed over four held-out graphs per size under the reduced exact-search protocol.

Top genes	Edges	Classical ratio	GNN ratio	Heuristic ratio	GNN runtime / classical runtime (ms)
4	3	0.9107	0.8619	0.8952	0.1459 / 136.8505
6	6	0.9214	0.9196	0.9189	0.1930 / 199.2276
8	11	0.8649	0.8699	0.8732	0.2806 / 286.0483
10	18	0.8676	0.8625	0.8667	0.4008 / 456.6395
12	26	0.8690	0.8694	0.8706	0.8261 / 1018.2954
14	36	0.8709	0.8697	0.8745	1.9952 / 2412.3906
16	48	0.8663	0.7935	0.7454	6.3308 / 8244.2970

[23] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, in *International Conference on Learning Representations* (2019).

[24] J. Gilmer, S. S. Schoenholz, P. Riley, O. Vinyals, and G. E. Dahl, in *International Conference on Machine Learning* (2017).

[25] K. Xu, C. Li, Y. Tian, T. Sonobe, K.-i. Kawarabayashi, and S. Jegelka, in *International Conference on Machine Learning* (2018).

[26] J. Klicpera, A. Bojchevski, and S. Gunnemann, in *International Conference on Learning Representations* (2019).

[27] Y. Rong, W. Huang, T. Xu, and J. Huang, in *International Conference on Learning Representations* (2020).

[28] M. Chen, Z. Wei, Z. Huang, B. Ding, and Y. Li, in *International Conference on Machine Learning* (2020).

[29] G. Corso, L. Cavalleri, D. Beaini, P. Li, and P. Velickovic, in *Advances in Neural Information Processing Systems* (2020).

[30] W. Hu, M. Fey, M. Zitnik, Y. Dong, H. Ren, B. Liu, M. Catasta, and J. Leskovec, in *Advances in Neural Information Processing Systems* (2020).

[31] V. P. Dwivedi, C. K. Joshi, T. Laurent, Y. Bengio, and X. Bresson, in *Advances in Neural Information Processing Systems* (2020).

[32] J. You, K. Gomes-Selman, R. Ying, and J. Leskovec, in *Advances in Neural Information Processing Systems* (2020).

[33] W. Hu, B. Liu, J. Gomes, M. Zitnik, P. Liang, V. Pande, and J. Leskovec, in *International Conference on Learning Representations* (2020).

[34] C. Morris, M. Ritzert, M. Fey, W. Hamilton, J. E. Lenssen, G. Rattan, and M. Grohe, in *AAAI Conference on Artificial Intelligence* (2019).

[35] H. Maron, H. Ben-Hamu, N. Shamir, and Y. Lipman, in *Advances in Neural Information Processing Systems* (2019).

[36] R. Ying, J. You, C. Morris, X. Ren, W. Hamilton, and J. Leskovec, in *Advances in Neural Information Processing Systems* (2018).

[37] H. Gao and S. Ji, in *International Conference on Machine Learning* (2019).

[38] J. Lee, I. Lee, and J. Kang, in *International Conference on Machine Learning* (2019).

[39] C. Ying, T. Cai, S. Luo, S. Zheng, G. Ke, D. He, Y. Shen, N. Lipka, and T.-Y. Liu, in *Advances in Neural Information Processing Systems* (2021).

[40] Q. Chen, Y. Song, C. Huang, H. Xia, Z. Wu, and X. Wang, in *Advances in Neural Information Processing Systems* (2022).

[41] T. Chen and C. Guestrin, in *ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (2016).

[42] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, in *Advances in Neural Information Processing Systems* (2017).

[43] L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin, in *Advances in Neural Information Processing Systems* (2018).

[44] S. O. Arik and T. Pfister, in *AAAI Conference on Artificial Intelligence* (2021).

[45] Y. Gorishniy, I. Rubachev, V. Khrulkov, and A. Babenko, in *Advances in Neural Information Processing Systems* (2021).

[46] R. Shwartz-Ziv and A. Armon, in *Advances in Neural Information Processing Systems* (2022).

[47] R. Caruana, Y. Lou, J. Gehrke, P. Koch, M. Sturm, and N. Elhadad, in *ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (2015).

[48] R. Miotto, L. Li, B. A. Kidd, and J. T. Dudley, in *ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (2016).

[49] E. Choi, A. Schuetz, W. F. Stewart, and J. Sun, in *KDD Workshop on Machine Learning for Healthcare* (2016).

[50] E. Choi, M. T. Bahadori, J. Sun, J. Kulas, A. Schuetz, and W. Stewart, in *ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (2017).

[51] S. Parisot, S. I. Ktena, E. Ferrante, M. Lee, R. Guerrero, B. Glocker, and D. Rueckert, Medical Image Analysis **48**, 117 (2018).

[52] A. Rajkomar, E. Oren, K. Chen, A. M. Dai, N. Hajaj, M. Hardt, P. J. Liu, X. Liu, J. Marcus, M. Sun, P. Sundberg, H. Yee, K. Zhang, Y. Zhang, G. Flores, G. E. Duggan, J. Irvine, Q. Le, K. Litsch, A. Mossin, J. Tansuwan, D. Wang, J. Wexler, J. Wilson, D. Ludwig, S. Volchenboum, K. Chou, M. Pearson, S. Madabushi, N. Shah, A. Butte, M. Howell, C. Cui, G. Corrado, J. Dean, and A. D. L. Scalable, npj Digital Medicine **1**, 18 (2018).

[53] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, in *International Conference on Machine Learning* (2017).

[54] Y. Gal and Z. Ghahramani, in *International Conference on Machine Learning* (2016).

[55] B. Lakshminarayanan, A. Pritzel, and C. Blundell, in *Advances in Neural Information Processing Systems* (2017).

[56] V. Kuleshov, N. Fenner, and S. Ermon, in *International Conference on Machine Learning* (2018).

[57] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. Dillon, B. Lakshminarayanan, and J. Snoek, in *Advances in Neural Information Processing*

Systems (2019).
[58] M. Minderer, J. Djolonga, R. Romijnders, F. Hubis, X. Zhai, N. Houlsby, S. Gelly, and M. Lucic, in *Advances in Neural Information Processing Systems* (2021).
[59] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollar, in *IEEE International Conference on Computer Vision* (2017).
[60] Y. Cui, M. Jia, T.-Y. Lin, Y. Song, and S. Belongie, in *IEEE Conference on Computer Vision and Pattern Recognition* (2019).
[61] J. Davis and M. Goadrich, in *International Conference on Machine Learning* (2006).
[62] H. Jiang, B. Kim, M. Y. Guan, and M. Gupta, in *Advances in Neural Information Processing Systems* (2018).
[63] C. Corbiere, N. Thome, A. Bar-Hen, M. Cord, and P. Perez, in *Advances in Neural Information Processing Systems* (2019).
[64] S. Thulasidasan, G. Chennupati, J. Bilmes, T. Bhattacharya, and S. Michalak, in *Advances in Neural Information Processing Systems* (2019).
[65] J. Wang, A. Ma, Y. Chang, J. Gong, Y. Jiang, R. Qi, C. Wang, H. Fu, D. Xu, and Q. Ma, Nature Communications **12**, 1882 (2021).
[66] L. Zhao and L. Akoglu, in *International Conference on Learning Representations* (2020).
[67] U. Alon and E. Yahav, in *International Conference on Learning Representations* (2021).
[68] J. Topping, F. Di Giovanni, B. Chamberlain, X. Dong, M. Bronstein, C. Twigg, and D. Beaini, in *International Conference on Learning Representations* (2022).
[69] K. K. Thekumparampil, C. Wang, S. Oh, and L. Li, in *ICLR Workshop* (2018).
[70] J. Klicpera, A. Bojchevski, and S. Gunnemann, in *NeurIPS Workshop on Graph Representation Learning* (2019).
[71] S. Brody, U. Alon, and E. Yahav, in *International Conference on Learning Representations* (2022).
[72] B. Knyazev, G. W. Taylor, and M. Amer, in *Advances in Neural Information Processing Systems* (2019).
[73] G. Li, M. Muller, A. K. Thabet, and B. Ghanem, in *IEEE International Conference on Computer Vision* (2019).
[74] J. Cai and S. Wang, in *ICML Workshop on Graph Representation Learning* (2020).
[75] V. P. Dwivedi and X. Bresson, in *AAAI Workshop on Deep Learning on Graphs* (2021).
[76] X. He *et al.*, in *NeurIPS Systems for Machine Learning Workshop* (2017).